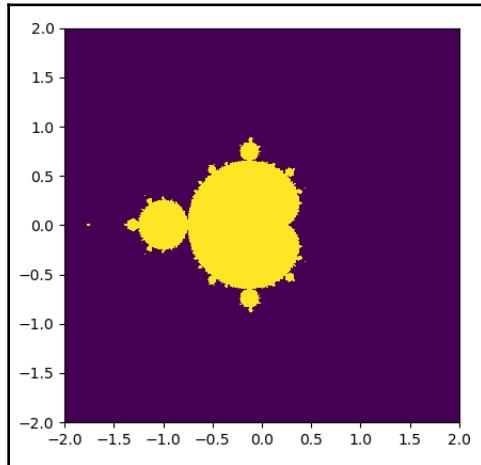


# Chapter 1: Why GPU Programming?



```
PS C:\Users\btuom\examples\1> python mandelbrot0.py
It took 14.617000103 seconds to calculate the Mandelbrot graph.
It took 0.110999822617 seconds to dump the image.
```

```
PS C:\Users\btuom\examples\1> python -m cProfile -s cumtime mandelbrot0.py > mandelbrot_profile.txt
PS C:\Users\btuom\examples\1>
```

```
It took 14.5690000057 seconds to calculate the Mandelbrot graph.
It took 0.136000156403 seconds to dump the image.
      564104 function calls (559254 primitive calls) in 14.965 seconds
```

Ordered by: cumulative time

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
1	0.002	0.002	14.966	14.966	mandelbrot0.py:1(<module>)
1	14.363	14.363	14.572	14.572	mandelbrot0.py:10(simple_mandelbrot)
263606	0.209	0.000	0.209	0.000	{range}
1	0.007	0.007	0.134	0.134	__init__.py:101(<module>)
1	0.003	0.003	0.123	0.123	pyplot.py:17(<module>)
12	0.017	0.001	0.119	0.010	__init__.py:1(<module>)
1	0.000	0.000	0.097	0.097	pyplot.py:694(savefig)
2	0.000	0.000	0.082	0.041	backend_agg.py:418(draw)
152/2	0.000	0.000	0.081	0.041	artist.py:47(draw_wrapper)
2	0.000	0.000	0.081	0.041	figure.py:1264(draw)
4/2	0.000	0.000	0.080	0.040	image.py:120(_draw_list_compositing_images)

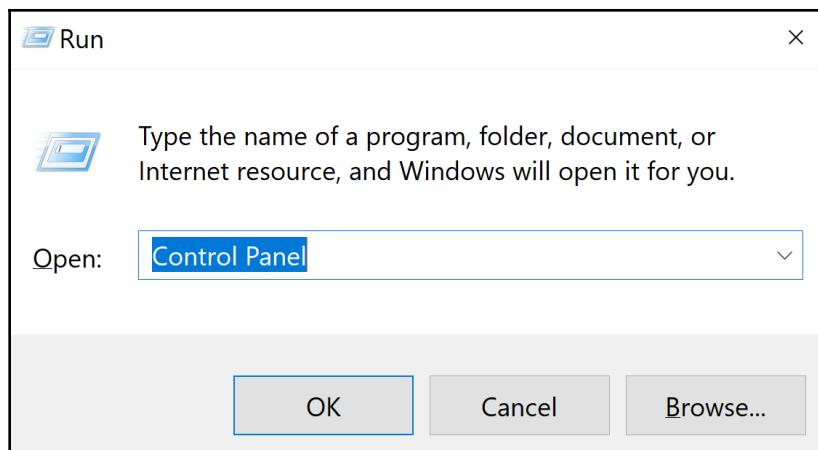
---

# Chapter 2: Setting Up Your GPU Programming Environment

```
Architecture:          x86_64
CPU op-mode(s):       32-bit, 64-bit
Byte Order:           Little Endian
CPU(s):               12
On-line CPU(s) list: 0-11
Thread(s) per core:  2
Core(s) per socket:  6
Socket(s):            1
NUMA node(s):         1
Vendor ID:            GenuineIntel
```

	total	used	free	shared	buff/cache	available
Mem:	15	3	9	0	2	12
Swap:	5	0	5			

```
01:00.0 VGA compatible controller: NVIDIA Corporation GP104M [GeForce GTX 1070 Mobile] (rev a1)
```



Control Panel\System and Security\System

Control Panel Home

Device Manager

Remote settings

System protection

Advanced system settings

View basic information about your computer

Windows edition

Windows 10 Pro

© 2018 Microsoft Corporation. All rights reserved.

Windows 10

System

Manufacturer: Microsoft Corporation Surface

Processor: Intel(R) Core(TM) i7-8650U CPU @ 1.90GHz 2.11 GHz

Installed memory (RAM): 8.00 GB

System type: 64-bit Operating System, x64-based processor

Pen and Touch: Pen and Touch Support with 10 Touch Points

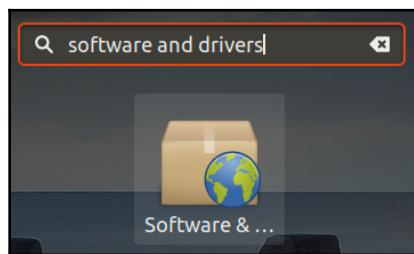
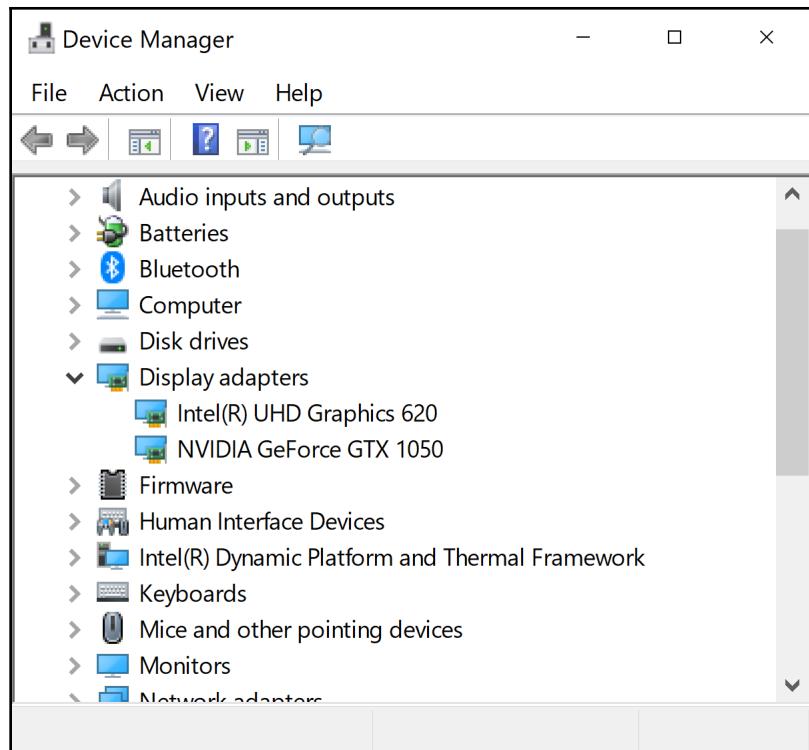
See also

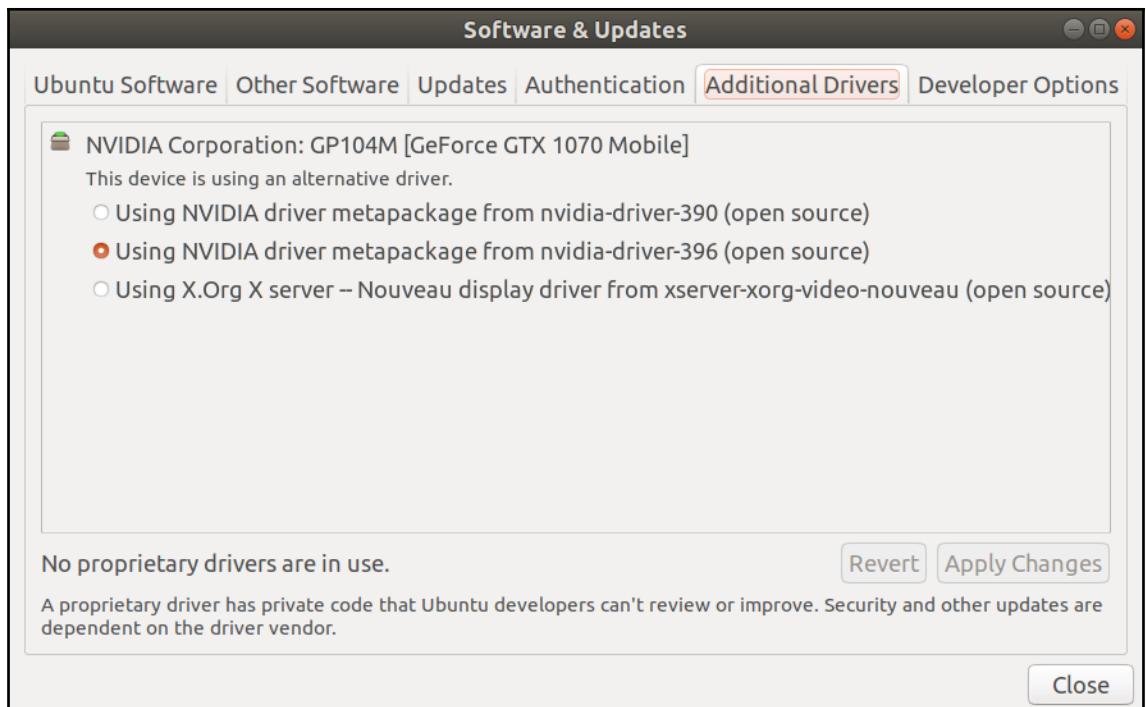
Microsoft Corporation support

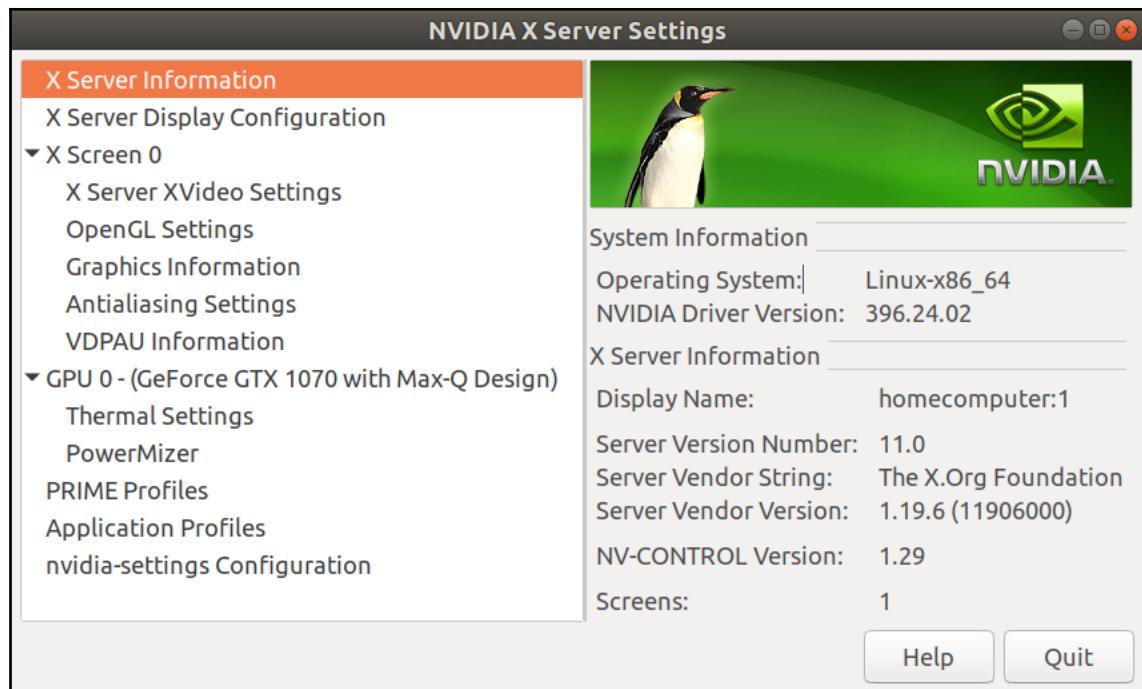
Security and Maintenance

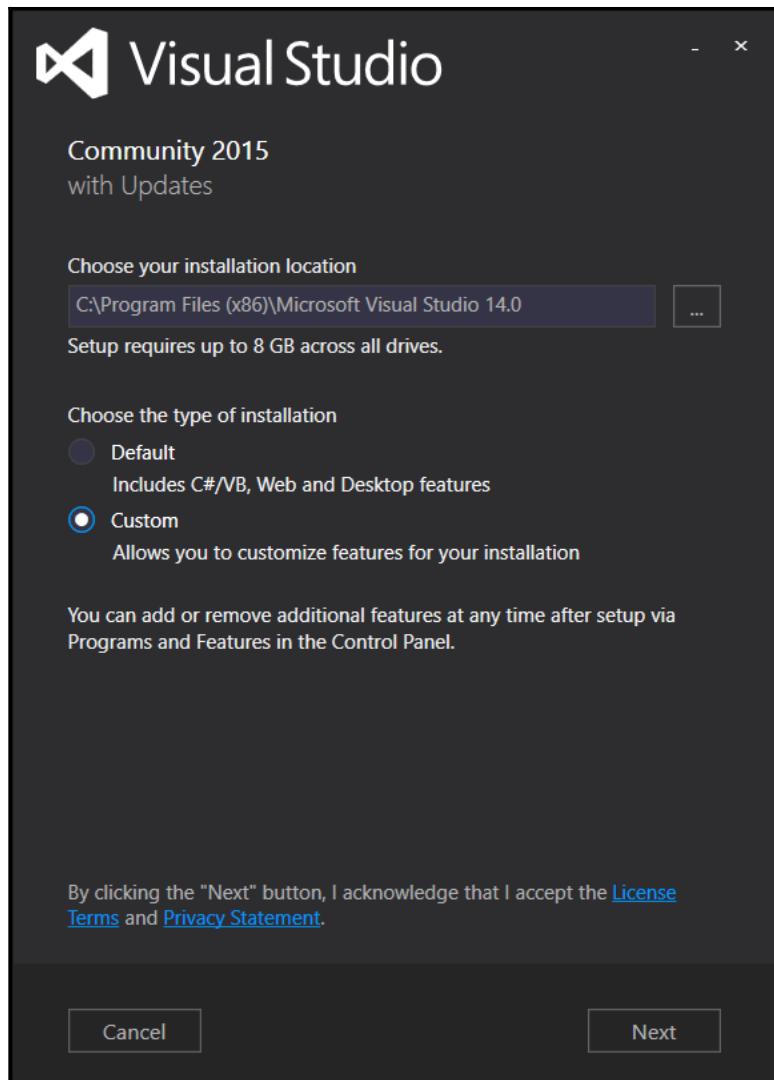
Website:

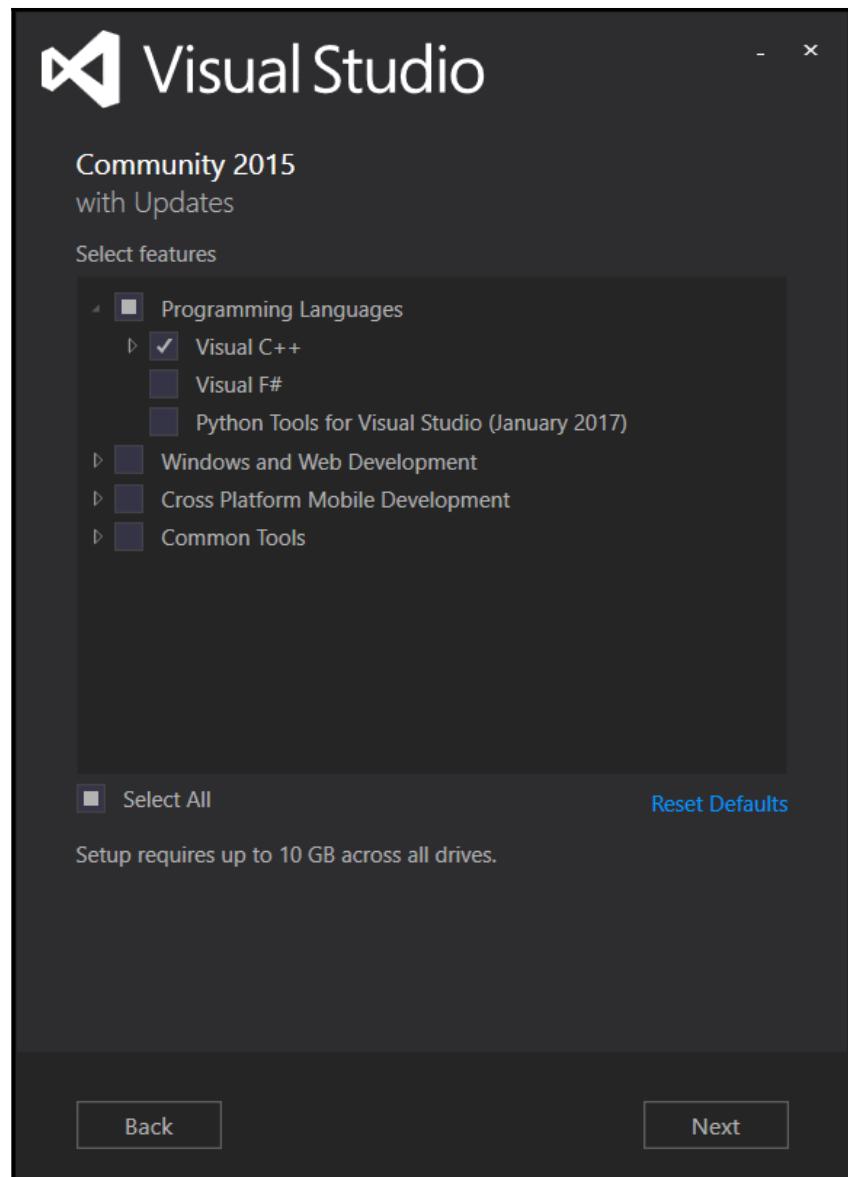
Online support











```
PS C:\Users\btuom\examples\3> python ./deviceQuery.py
CUDA device query (PyCUDA version)

Detected 1 CUDA Capable device(s)

Device 0: GeForce GTX 1050
```

# Chapter 3: Getting Started with PyCUDA

```
PS C:\ProgramData\NVIDIA Corporation\CUDA Samples\v9.1\bin\win64\Debug> .\deviceQuery.exe
C:\ProgramData\NVIDIA Corporation\CUDA Samples\v9.1\bin\win64\Debug\deviceQuery.exe Starting...

CUDA Device Query (Runtime API) version (CUDART static linking)

Detected 1 CUDA Capable device(s)

Device 0: "GeForce GTX 1050"
  CUDA Driver Version / Runtime Version      9.1 / 9.1
  CUDA Capability Major/Minor version number: 6.1
  Total amount of global memory:            2048 MBytes (2147483648 bytes)
    ( 5) Multiprocessors, (128) CUDA Cores/MP:
  GPU Max Clock rate:                     1493 MHz (1.49 GHz)
  Memory Clock rate:                      3504 Mhz
  Memory Bus Width:                       128-bit
  L2 Cache Size:                          524288 bytes
  Maximum Texture Dimension Size (x,y,z)   1D=(131072), 2D=(131072, 65536), 3D=(16384, 16384, 16384)
  Maximum Layered 1D Texture Size, (num) layers 1D=(32768), 2048 layers
  Maximum Layered 2D Texture Size, (num) layers 2D=(32768, 32768), 2048 layers
  Total amount of constant memory:          65536 bytes
  Total amount of shared memory per block:  49152 bytes
  Total number of registers available per block: 65536
  Warp size:                            32
  Maximum number of threads per multiprocessor: 2048
  Maximum number of threads per block:       1024
  Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
  Max dimension size of a grid size (x,y,z): (2147483647, 65535, 65535)
  Maximum memory pitch:                  2147483647 bytes
  Texture alignment:                    512 bytes
  Concurrent copy and kernel execution: Yes with 2 copy engine(s)
  Run time limit on kernels:           No
  Integrated GPU sharing Host Memory: No
  Support host page-locked memory mapping: Yes
  Alignment requirement for Surfaces: Yes
  Device has ECC support:             Disabled
  CUDA Device Driver Mode (TCC or WDDM): WDDM (Windows Display Driver Model)
  Device supports Unified Addressing (UVA): Yes
  Supports Cooperative Kernel Launch:  No
  Supports MultiDevice Co-op Kernel Launch: No
  Device PCI Domain ID / Bus ID / location ID: 0 / 2 / 0
  Compute Mode:
    < Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >

deviceQuery, CUDA Driver = CUDART, CUDA Driver Version = 9.1, CUDA Runtime Version = 9.1, NumDevs = 1
Result = PASS
```

```
In [8]: import pycuda.driver as drv
```

```
In [9]: drv.init()
```

```
In [10]: print 'Detected {} CUDA Capable device(s)'.format(drv.Device.count())
Detected 1 CUDA Capable device(s)
```

---

```
PS C:\Users\btruom\examples\3> python deviceQuery.py
CUDA device query (PyCUDA version)

Detected 1 CUDA Capable device(s)

Device 0: GeForce GTX 1050
    Compute Capability: 6.1
    Total Memory: 2048 megabytes
    (5) Multiprocessors, (128) CUDA Cores / Multiprocessor: 640 CUDA Cores
    MAXIMUM_TEXTURE2D_LINEAR_PITCH: 2097120
    MAXIMUM_TEXTURE2D_GATHER_WIDTH: 32768
    MAXIMUM_TEXTURE2D_GATHER_HEIGHT: 32768
    PCI_DEVICE_ID: 0
    MAXIMUM_TEXTURE3D_WIDTH: 16384
    MAXIMUM_SURFACE2D_WIDTH: 131072
    MAXIMUM_TEXTURE1D_MIPMAPPED_WIDTH: 16384
    GLOBAL_MEMORY_BUS_WIDTH: 128
    LOCAL_L1_CACHE_SUPPORTED: 1
    MAXIMUM_SURFACE3D_DEPTH: 16384
    MAXIMUM_TEXTURE3D_HEIGHT: 16384
    PCI_DOMAIN_ID: 0
    COMPUTE_CAPABILITY_MINOR: 1
    MULTI_GPU_BOARD_GROUP_ID: 0
    MAX_REGISTERS_PER_BLOCK: 65536
    MAXIMUM_TEXTURE2D_ARRAY_WIDTH: 32768
    COMPUTE_CAPABILITY_MAJOR: 6
    MAXIMUM_SURFACE2D_LAYERED_HEIGHT: 32768
    MAXIMUM_TEXTURE1D_LAYERED_LAYERS: 2048
    UNIFIED_ADDRESSING: 1
```

---

```
In [24]: import numpy as np
In [25]: import pycuda.autoinit
In [26]: from pycuda import gpuarray
In [27]: host_data = np.array([1,2,3,4,5],dtype=np.float32)
In [28]: device_data = gpuarray.to_gpu(host_data)
In [29]: device_data_x2 = 2 * device_data
In [30]: host_data_x2 = device_data_x2.get()
In [31]: print host_data_x2
[ 2.  4.  6.  8. 10.]
In [32]:
```

```
In [14]: x_host = np.array([1,2,3], dtype=np.float32)
In [15]: y_host = np.array([1,1,1], dtype=np.float32)
In [16]: z_host = np.array([2,2,2], dtype=np.float32)
In [17]: x_device = gpuarray.to_gpu(x_host)
In [18]: y_device = gpuarray.to_gpu(y_host)
In [19]: z_device = gpuarray.to_gpu(z_host)
In [20]: x_host + y_host
Out[20]: array([ 2.,  3.,  4.], dtype=float32)
In [21]: (x_device + y_device).get()
Out[21]: array([ 2.,  3.,  4.], dtype=float32)
In [22]: x_host ** z_host
Out[22]: array([ 1.,  4.,  9.], dtype=float32)
In [23]: (x_device ** z_device).get()
Out[23]: array([ 1.,  4.,  9.], dtype=float32)
In [24]: x_host / x_host
Out[24]: array([ 1.,  1.,  1.], dtype=float32)
In [25]: (x_device / x_device).get()
Out[25]: array([ 1.,  1.,  1.], dtype=float32)
In [26]: z_host - x_host
Out[26]: array([ 1.,  0., -1.], dtype=float32)
In [27]: (z_device - x_device).get()
Out[27]: array([ 1.,  0., -1.], dtype=float32)
In [28]: z_host / 2
Out[28]: array([ 1.,  1.,  1.], dtype=float32)
In [29]: (z_device / 2).get()
Out[29]: array([ 1.,  1.,  1.], dtype=float32)
In [30]: x_host - 1
Out[30]: array([ 0.,  1.,  2.], dtype=float32)
In [31]: (x_device - 1).get()
Out[31]: array([ 0.,  1.,  2.], dtype=float32)
```

---

```
In [1]: run time_calc0.py
total time to compute on CPU: 0.078000
total time to compute on GPU: 1.094000
Is the host computation the same as the GPU computation? : True

In [2]: run time_calc0.py
total time to compute on CPU: 0.079000
total time to compute on GPU: 0.008000
Is the host computation the same as the GPU computation? : True

In [3]: run time_calc0.py
total time to compute on CPU: 0.080000
total time to compute on GPU: 0.007000
Is the host computation the same as the GPU computation? : True

In [4]: run time_calc0.py
total time to compute on CPU: 0.078000
total time to compute on GPU: 0.009000
Is the host computation the same as the GPU computation? : True

In [5]: run time_calc0.py
total time to compute on CPU: 0.079000
total time to compute on GPU: 0.009000
Is the host computation the same as the GPU computation? : True
```

---

```
In [2]: %prun -s cumulative exec(time_calc_code)
total time to compute on CPU: 0.078000
total time to compute on GPU: 1.100000
Is the host computation the same as the GPU computation? : True
    17353 function calls (17146 primitive calls) in 3.175 seconds

Ordered by: cumulative time

      ncalls  tottime  percall  cumtime  percall filename:lineno(function)
          1    0.000    0.000    1.101    1.101 gpuarray.py:452(__mul__)
          1    0.000    0.000    1.092    1.092 gpuarray.py:317(_axpbz)
          1    0.000    0.000    1.091    1.091 <decorator-gen-122>:1(get_axpbz_kernel)

          1    0.000    0.000    1.091    1.091 tools.py:414(context_dependent_memoize)

          1    0.000    0.000    1.091    1.091 elementwise.py:413(get_axpbz_kernel)
          1    0.000    0.000    1.091    1.091 elementwise.py:155(get_elwise_kernel)
          1    0.000    0.000    1.091    1.091 elementwise.py:126(get_elwise_kernel_an
d_types)
          1    0.000    0.000    1.091    1.091 elementwise.py:41(get_elwise_module)
          1    0.001    0.001    1.089    1.089 compiler.py:285(__init__)
          1    0.001    0.001    1.089    1.089 compiler.py:190(compile)
          1    0.001    0.001    1.070    1.070 compiler.py:69(compile_plain)
          2    0.000    0.000    1.061    0.531 prefork.py:222(call_capture_output)
          2    0.000    0.000    1.061    0.531 prefork.py:43(call_capture_output)
          1    0.000    0.000    0.950    0.950 compiler.py:36(preprocess_source)
          2    0.000    0.000    0.837    0.419 subprocess.py:448(communicate)
          2    0.000    0.000    0.837    0.419 subprocess.py:698(__communicate)
          6    0.000    0.000    0.836    0.139 threading.py:309(wait)
```

---

```
In [3]: %prun -s cumulative exec(time_calc_code)
total time to compute on CPU: 0.101000
total time to compute on GPU: 0.015000
Is the host computation the same as the GPU computation? : True
    342 function calls (336 primitive calls) in 1.315 seconds

Ordered by: cumulative time

      ncalls  tottime  percall  cumtime  percall filename:lineno(function)
          1    0.000    0.000    1.606    1.606 <string>:1(<module>)
          1    0.016    0.016    0.650    0.650 numeric.py:2397(allclose)
          1    0.069    0.069    0.630    0.630 numeric.py:2463(isclose)
          1    0.400    0.400    0.554    0.554 numeric.py:2522(within_tol)
          1    0.452    0.452    0.452    0.452 {method 'random_sample' of 'mtrand.Ran
omState' objects}
          2    0.191    0.096    0.191    0.096 gpuarray.py:1174(_memcpy_discontig)
          2    0.154    0.077    0.154    0.077 {abs}
          1    0.000    0.000    0.107    0.107 gpuarray.py:248(get)
          1    0.000    0.000    0.094    0.094 gpuarray.py:990(to_gpu)
          1    0.000    0.000    0.085    0.085 gpuarray.py:230(set)
          2    0.018    0.009    0.018    0.009 gpuarray.py:162(__init__)
          3    0.000    0.000    0.012    0.004 fromnumeric.py:1973(all)
```

```
PS C:\Users\btuom\examples\3> python simple_element_kernel_example0.py
total time to compute on CPU: 0.092000
total time to compute on GPU: 1.494000
Is the host computation the same as the GPU computation? : True
PS C:\Users\btuom\examples\3>
```

---

```
In [1]: run simple_element_kernel_example0.py
total time to compute on CPU: 0.080000
total time to compute on GPU: 0.989000
Is the host computation the same as the GPU computation? : True

In [2]: speedcomparison()
total time to compute on CPU: 0.081000
total time to compute on GPU: 0.000000
Is the host computation the same as the GPU computation? : True

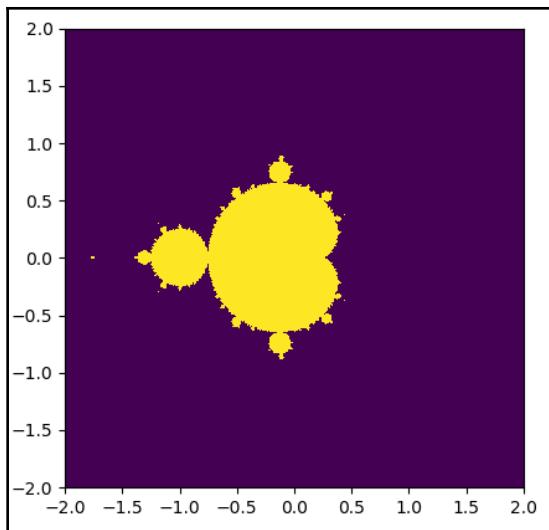
In [3]: speedcomparison()
total time to compute on CPU: 0.096000
total time to compute on GPU: 0.000000
Is the host computation the same as the GPU computation? : True

In [4]: speedcomparison()
total time to compute on CPU: 0.085000
total time to compute on GPU: 0.000000
Is the host computation the same as the GPU computation? : True

In [5]: speedcomparison()
total time to compute on CPU: 0.085000
total time to compute on GPU: 0.000000
Is the host computation the same as the GPU computation? : True

In [6]:
```

```
In [1]: run gpu_mandelbrot0.py
It took 0.894000053406 seconds to calculate the Mandelbrot graph.
It took 0.102999925613 seconds to dump the image.
```



```
In [2]: pow2 = lambda x : x**2
```

```
In [3]: pow2(2)  
Out[3]: 4
```

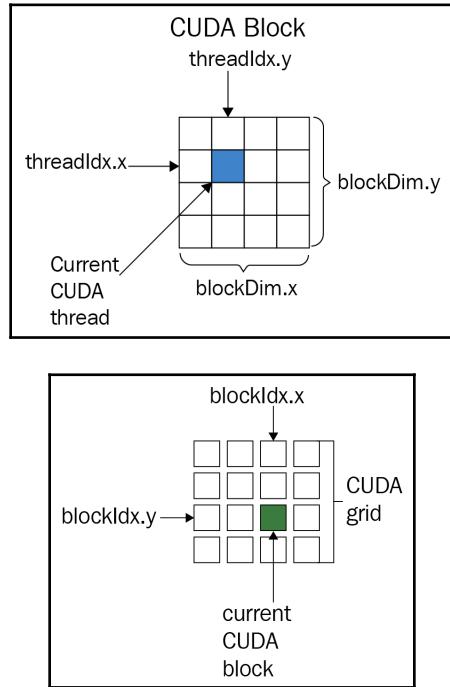
```
In [4]: pow2(3)  
Out[4]: 9
```

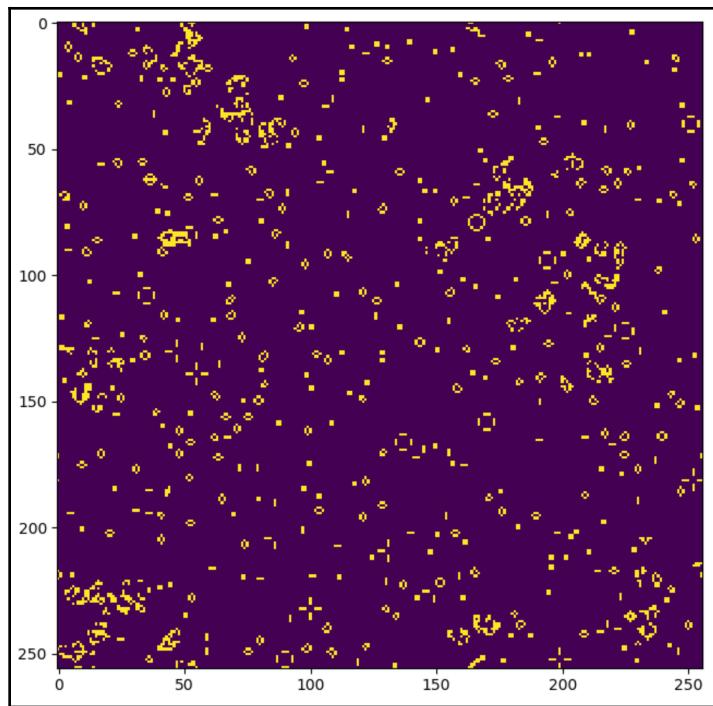
```
In [5]: pow2(4)  
Out[5]: 16
```

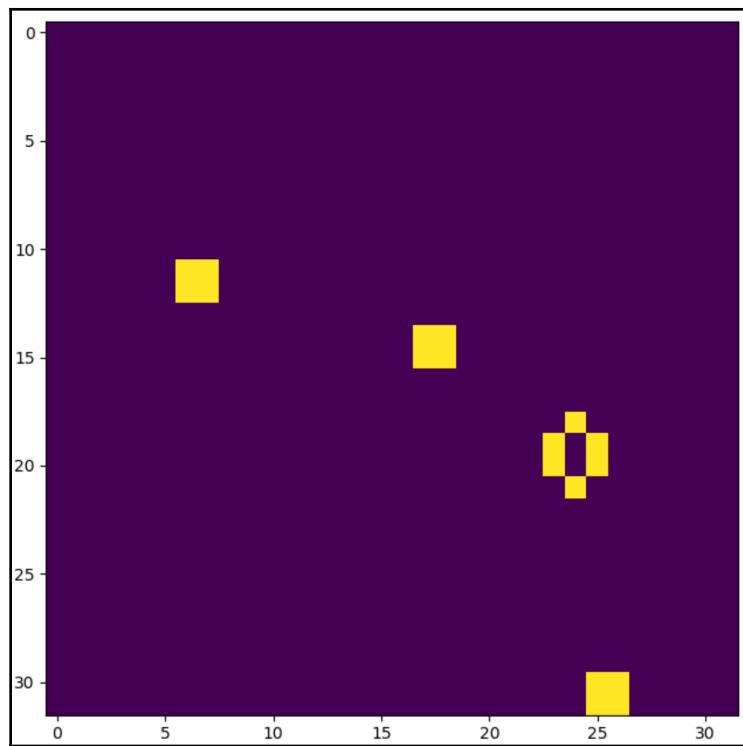
```
In [6]: map(lambda x : x**2, [2,3,4])  
Out[6]: [4, 9, 16]
```

```
In [1]: run simple_scankernel0.py  
[ 1  3  6 10]  
[ 1  3  6 10]
```

# Chapter 4: Kernels, Threads, Blocks, and Grids



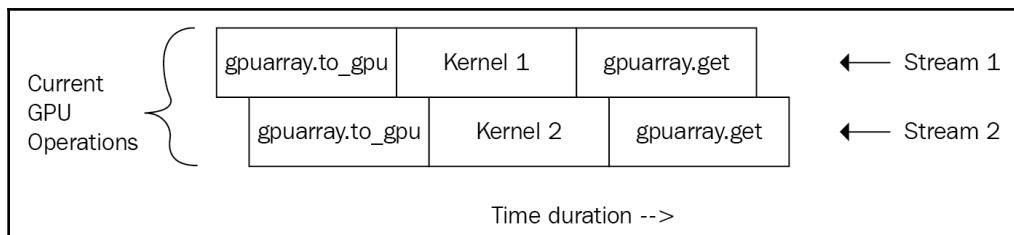
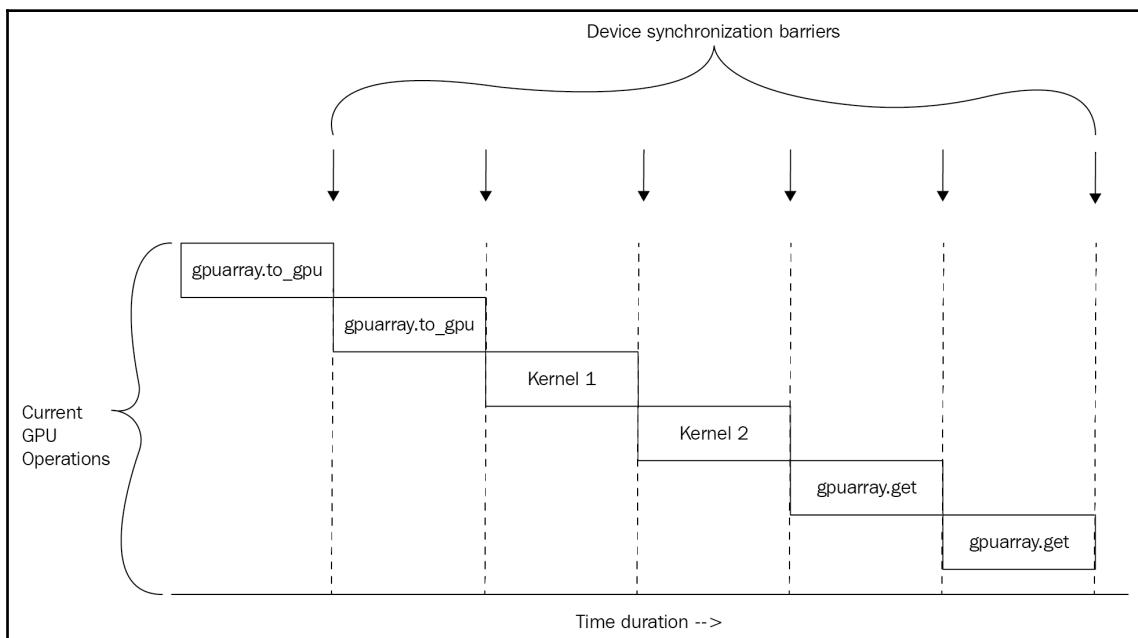


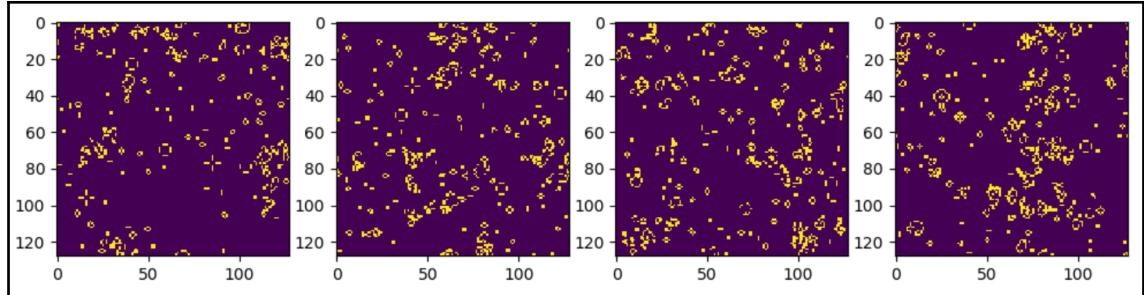


# Chapter 5: Streams, Events, Contexts, and Concurrency

```
PS C:\Users\btuom\examples\5> python .\multi-kernel.py
Total time: 2.976000
```

```
PS C:\Users\btuom\examples\5> python .\multi-kernel_streams.py
Total time: 0.945000
```





```
PS C:\Users\btruom\examples\5> python .\simple_event_example.py
Has the kernel started yet? False
Has the kernel ended yet? False
```

```
PS C:\Users\btruom\examples\5> python .\simple_event_example.py
Has the kernel started yet? True
Has the kernel ended yet? True
Kernel execution time in milliseconds: 1047.391235
```

```
PS C:\Users\btruom\examples\5> python .\multi-kernel_events.py
Total time: 1.078000
Mean kernel duration (milliseconds): 71.417903
Mean kernel standard deviation (milliseconds): 6.401030
```

```
PS C:\Users\btruom\examples\5> python .\single_thread_example.py
Hello from the thread you just spawned!
The thread completed and returned this value: 123
```

---

# Chapter 6: Debugging and Profiling Your CUDA Code

```
PS C:\Users\btuom\examples\6> python .\hello-world_gpu.py
Hello world from thread 0, in block 1!
Hello world from thread 1, in block 1!
Hello world from thread 2, in block 1!
Hello world from thread 3, in block 1!
Hello world from thread 4, in block 1!
Hello world from thread 0, in block 0!
Hello world from thread 1, in block 0!
Hello world from thread 2, in block 0!
Hello world from thread 3, in block 0!
Hello world from thread 4, in block 0!
-----
This kernel was launched over a grid consisting of 2 blocks,
where each block has 5 threads.
```

```
PS C:\Users\btuom\examples\6> python .\broken_matrix_ker.py
Traceback (most recent call last):
  File ".\broken_matrix_ker.py", line 64, in <module>
    assert( np.allclose(output_mat_gpu.get(), output_mat) )
AssertionError
PS C:\Users\btuom\examples\6>
```

```
PS C:\Users\btuom\examples\6> python .\broken_matrix_ker.py
threadIdx.x,y: 0,0 blockIdx.x,y: 1,0 -- row is 1, col is 0.
threadIdx.x,y: 1,0 blockIdx.x,y: 1,0 -- row is 2, col is 0.
threadIdx.x,y: 0,1 blockIdx.x,y: 1,0 -- row is 1, col is 1.
threadIdx.x,y: 1,1 blockIdx.x,y: 1,0 -- row is 2, col is 1.
threadIdx.x,y: 0,0 blockIdx.x,y: 1,1 -- row is 1, col is 1.
threadIdx.x,y: 1,0 blockIdx.x,y: 1,1 -- row is 2, col is 1.
threadIdx.x,y: 0,1 blockIdx.x,y: 1,1 -- row is 1, col is 2.
threadIdx.x,y: 1,1 blockIdx.x,y: 1,1 -- row is 2, col is 2.
threadIdx.x,y: 0,0 blockIdx.x,y: 0,0 -- row is 0, col is 0.
threadIdx.x,y: 1,0 blockIdx.x,y: 0,0 -- row is 1, col is 0.
threadIdx.x,y: 0,1 blockIdx.x,y: 0,0 -- row is 0, col is 1.
threadIdx.x,y: 1,1 blockIdx.x,y: 0,0 -- row is 1, col is 1.
threadIdx.x,y: 0,0 blockIdx.x,y: 0,1 -- row is 0, col is 1.
threadIdx.x,y: 1,0 blockIdx.x,y: 0,1 -- row is 1, col is 1.
threadIdx.x,y: 0,1 blockIdx.x,y: 0,1 -- row is 0, col is 2.
threadIdx.x,y: 1,1 blockIdx.x,y: 0,1 -- row is 1, col is 2.
Traceback (most recent call last):
  File ".\broken_matrix_ker.py", line 64, in <module>
    assert( np.allclose(output_mat_gpu.get(), output_mat) )
AssertionError
PS C:\Users\btuom\examples\6>
```

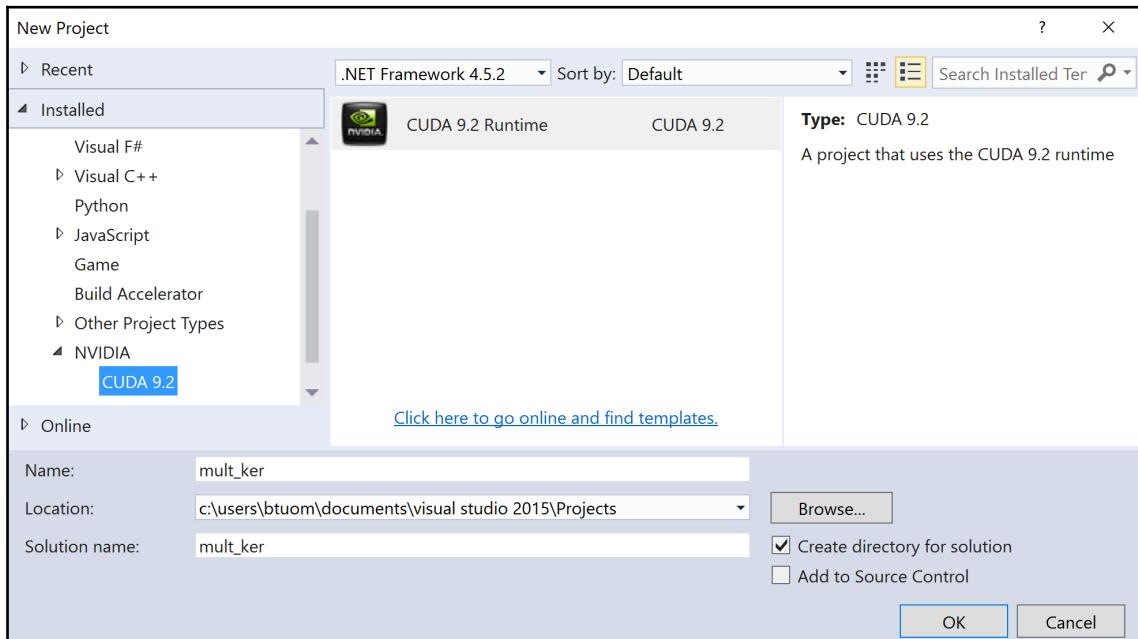
```
In [2]: print test_a
[[ 1.  2.  3.  4.]
 [ 1.  2.  3.  4.]
 [ 1.  2.  3.  4.]
 [ 1.  2.  3.  4.]]

In [3]: print test_b
[[ 14.  13.  12.  11.]
 [ 14.  13.  12.  11.]
 [ 14.  13.  12.  11.]
 [ 14.  13.  12.  11.]]
```

In [4]:

```
Dot-product loop: k value is 0, matrix_a value is 1.000000, matrix_b is 14.000000.
Dot-product loop: k value is 1, matrix_a value is 1.000000, matrix_b is 13.000000.
Dot-product loop: k value is 2, matrix_a value is 1.000000, matrix_b is 12.000000.
Dot-product loop: k value is 3, matrix_a value is 1.000000, matrix_b is 11.000000.
```

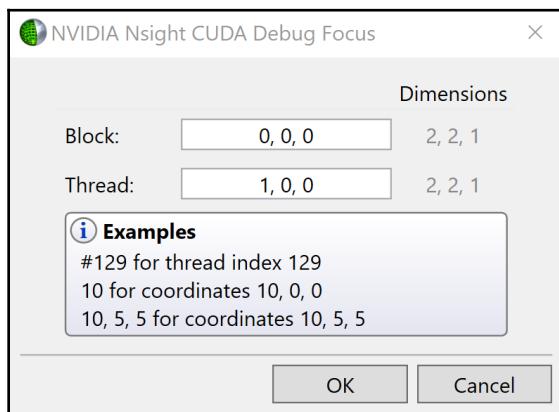
```
PS C:\Users\btruom\examples\6> nvcc matrix_ker.cu -o matrix_ker
matrix_ker.cu
    Creating library matrix_ker.lib and object matrix_ker.exp
PS C:\Users\btruom\examples\6> ./matrix_ker.exe
Success! Output of kernel matches expected output.
PS C:\Users\btruom\examples\6>
```



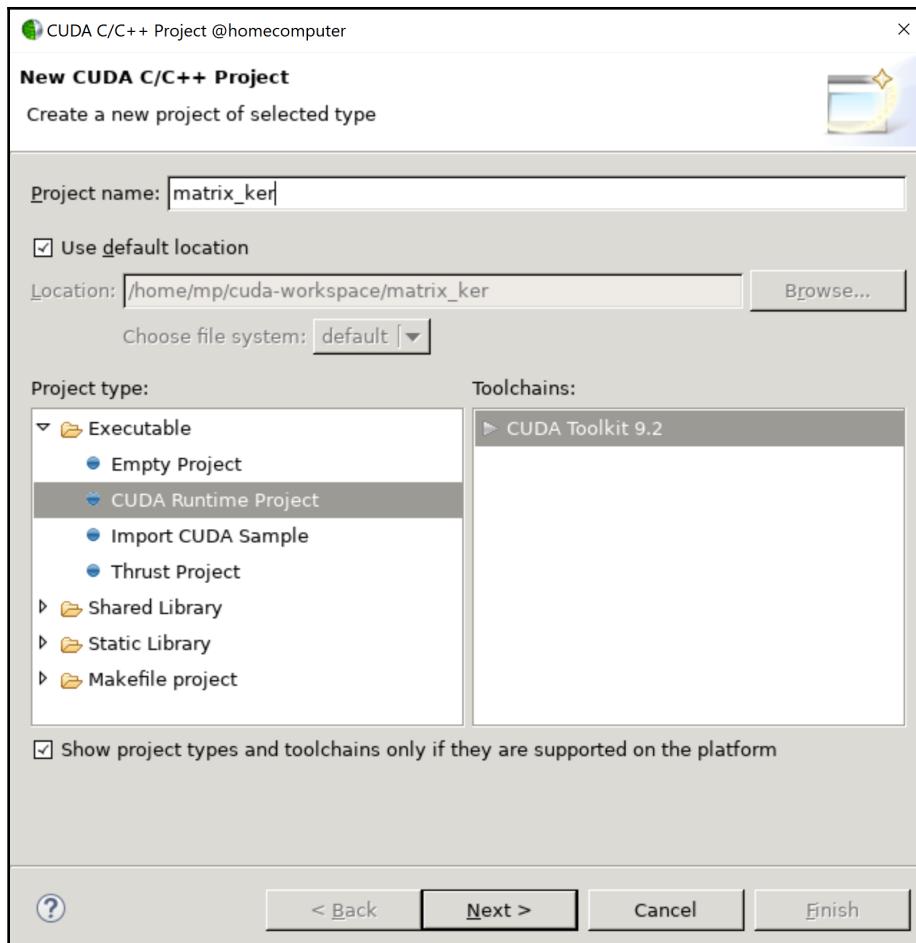
The screenshot shows the code editor with the file 'matrix\_ker.cu' open. The code implements a parallel matrix multiplication kernel using CUDA. It defines a global function 'matrix\_mult\_ker' that takes pointers to matrices 'matrix\_a' and 'matrix\_b', and a pointer to the output matrix 'output\_matrix', along with the dimension 'N'. The function iterates over row and column indices, calculating the sum of products using the 'rowcol\_dot' function. A red circle highlights the first line of the function definition. The code editor interface includes tabs for other files like 'matrix.h' and 'matrix.h', a status bar showing '100 %', and a bottom navigation bar.

```
37
38     // matrix multiplication kernel that is parallelized over row/column tuples.
39     __global__ void matrix_mult_ker(float * matrix_a, float * matrix_b, float * output_matrix, int N)
40     {
41
42         int row = blockIdx.x*blockDim.x + threadIdx.x;
43         int col = blockIdx.y*blockDim.y + threadIdx.y;
44
45         output_matrix[col + row*N] = rowcol_dot(matrix_a, matrix_b, row, col, N);
46     }
47
```

Locals		
Name	Value	Type
@flatBlockIdx	0	long
@flatThreadIdx	0	long
threadIdx	{x = 0, y = 0, z = 0}	const uint
x	0	unsigned
y	0	unsigned
z	0	unsigned
blockIdx	{x = 0, y = 0, z = 0}	const uint
x	0	unsigned
y	0	unsigned
z	0	unsigned
blockDim	{x = 2, y = 2, z = 1}	const dim3
gridDim	{x = 2, y = 2, z = 1}	const dim3
@gridId	1	const long
row	0	int
col	'col' has no value at the target location.	
matrix	0x0000000502400000	device



Locals		
Name	Value	Type
@flatBlockIdx	0	long
@flatThreadIdx	1	long
threadIdx	{x = 1, y = 0, z = 0}	const uin
x	1	unsigned
y	0	unsigned
z	0	unsigned
blockIdx	{x = 0, y = 0, z = 0}	const uin
x	0	unsigned
y	0	unsigned
z	0	unsigned
blockDim	{x = 2, y = 2, z = 1}	const din
gridDim	{x = 2, y = 2, z = 1}	const din
@gridId	1	const lon
row	1	int
col	'col' has no value at the target location.	
matrix	0x0000000502400000_1	device



The screenshot shows the Eclipse IDE's 'Console' view. The title bar includes tabs for 'Problems', 'Tasks', 'Console', 'Properties', and other icons. The main area displays the output of a terminated process:

```
<terminated> matrix_ker [C/C++ Application] /home/mp/cuda-workspace/matrix_ker/Debug/matrix_ker (8/19/18 5:07 PM)
Success! Output of kernel matches expected output.
```

(x)= Variables Breakpoints Registers CUDA Modules

Search CUDA Information

[0] matrix_mult_ker	Device 0 (GP104-A)	4 blocks of 4 are running
↳ (0,0,0)	SM 0	4 threads of 4 are running
↳ (0,0,0)	Warp 0 Lane 0	matrix_ker.cu:43 (0x555555ca7910)
↳ (0,1,0)	Warp 0 Lane 2	matrix_ker.cu:43 (0x555555ca7910)
↳ (1,0,0)	Warp 0 Lane 1	matrix_ker.cu:43 (0x555555ca7910)
↳ (1,1,0)	Warp 0 Lane 3	matrix_ker.cu:43 (0x555555ca7910)
↳ (0,1,0)	SM 2	4 threads of 4 are running
↳ (1,0,0)	SM 1	4 threads of 4 are running
↳ (1,1,0)	SM 3	4 threads of 4 are running

(x)= Variables Breakpoints Registers CUDA Modules

Name Type Value

Name	Type	T(0,0,0)B(0,0,0)
↳ matrix_a	@generic float * @parameter	0x7fffca400000
↳ matrix_b	@generic float * @parameter	0x7fffca400200
↳ output_matrix	@generic float * @parameter	0x7fffca400400
(x)= N	@parameter int	4
(x)= row	@register int	0
(x)= col	@register int	<optimized out>

(x)= Variables Breakpoints Registers CUDA Modules

Name Type Value

Name	Type	Value
↳ matrix_a	@generic float * @parameter	0x7fffca400000
↳ matrix_b	@generic float * @parameter	0x7fffca400200
↳ output_matrix	@generic float * @parameter	0x7fffca400400
(x)= N	@parameter int	4
(x)= row	@register int	1
(x)= col	@register int	<optimized out>

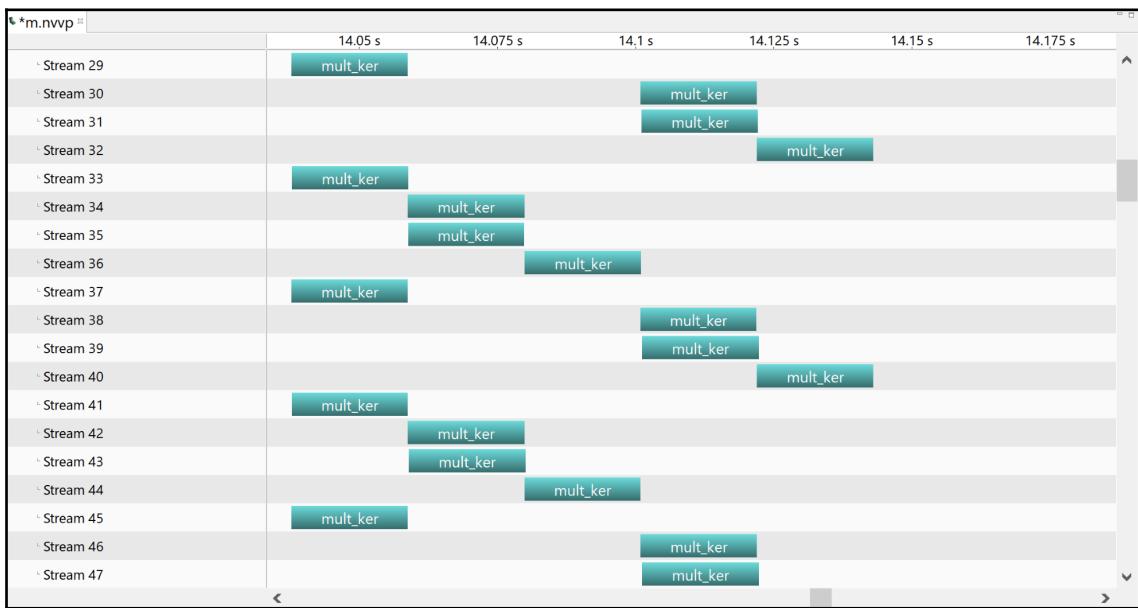
```
PS C:\Users\btuom\examples\6> .\divergence_test.exe
threadIdx.x 0 : This is an even thread.
threadIdx.x 2 : This is an even thread.
threadIdx.x 4 : This is an even thread.
threadIdx.x 6 : This is an even thread.
threadIdx.x 8 : This is an even thread.
threadIdx.x 10 : This is an even thread.
threadIdx.x 12 : This is an even thread.
threadIdx.x 14 : This is an even thread.
threadIdx.x 16 : This is an even thread.
threadIdx.x 18 : This is an even thread.
threadIdx.x 20 : This is an even thread.
threadIdx.x 22 : This is an even thread.
threadIdx.x 24 : This is an even thread.
threadIdx.x 26 : This is an even thread.
threadIdx.x 28 : This is an even thread.
threadIdx.x 30 : This is an even thread.
threadIdx.x 1 : This is an odd thread.
threadIdx.x 3 : This is an odd thread.
threadIdx.x 5 : This is an odd thread.
threadIdx.x 7 : This is an odd thread.
threadIdx.x 9 : This is an odd thread.
threadIdx.x 11 : This is an odd thread.
threadIdx.x 13 : This is an odd thread.
threadIdx.x 15 : This is an odd thread.
threadIdx.x 17 : This is an odd thread.
threadIdx.x 19 : This is an odd thread.
threadIdx.x 21 : This is an odd thread.
threadIdx.x 23 : This is an odd thread.
threadIdx.x 25 : This is an odd thread.
threadIdx.x 27 : This is an odd thread.
threadIdx.x 29 : This is an odd thread.
threadIdx.x 31 : This is an odd thread.
PS C:\Users\btuom\examples\6>
```

The screenshot shows a debugger interface with two main panes. The left pane displays the CUDA source code for `divergence_ker.cu`. The right pane shows the CUDA memory dump for Device 0 (GP104-A), which includes a tree view of the memory hierarchy and a detailed table for each lane.

**CUDA Memory Dump:**

Lane	Address	Value	Type	Description
(0,0,0)	(0,0,0)	Device 0 (GP104-A)	Device	1 block of 1 is running
(0,0,0)	(0,0,0)	SM 0	SM	32 threads of 32
(0,0,0)	(0,0,0)	Warp 0 Lane 0	Warp	divergence_ker.c
(1,0,0)	(1,0,0)	Warp 0 Lane 1	Warp	divergence_ker.c
(2,0,0)	(2,0,0)	Warp 0 Lane 2	Warp	divergence_ker.c
(3,0,0)	(3,0,0)	Warp 0 Lane 3	Warp	divergence_ker.c

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	42.94%	2.3360us	2	1.1680us	896ns	1.4400us	[CUDA memcpy HtoD]
	41.18%	2.2400us	1	2.2400us	2.2400us	2.2400us	matrix_mult_ker(float*, float*, float*, int)
	15.88%	864ns	1	864ns	864ns	864ns	[CUDA memcpy DtoH]
API calls:	72.42%	139.27ms	3	46.422ms	7.7580us	139.25ms	cudaMalloc
	25.66%	49.351ms	1	49.351ms	49.351ms	49.351ms	cudaDeviceReset
	1.46%	2.8053ms	88	31.878us	484ns	1.5375ms	cuDeviceGetAttribute
	0.15%	290.91us	3	96.969us	14.060us	260.85us	cudaFree
	0.14%	266.18us	3	88.727us	73.212us	111.52us	cudaMemcpy
	0.06%	119.27us	1	119.27us	119.27us	119.27us	cuDeviceGetName
	0.05%	101.33us	2	50.666us	11.152us	90.181us	cudaDeviceSynchronize
	0.02%	38.787us	1	38.787us	38.787us	38.787us	cuDeviceTotalMem
	0.02%	29.576us	1	29.576us	29.576us	29.576us	cudaLaunchKernel
	0.01%	21.818us	1	21.818us	21.818us	21.818us	cudaSetDevice
	0.00%	8.7280us	3	2.9090us	485ns	7.2730us	cuDeviceGetCount
	0.00%	8.7270us	1	8.7270us	8.7270us	8.7270us	cuDeviceGetPCIBusId
	0.00%	3.3940us	2	1.6970us	485ns	2.9090us	cuDeviceGet



# Chapter 7: Using the CUDA Libraries with Scikit-CUDA

```
In [3]: run cublas_gemm_flops.py
Single-precision performance: 1748.4264918 GFLOPS
Double-precision performance: 61.7956005349 GFLOPS
```



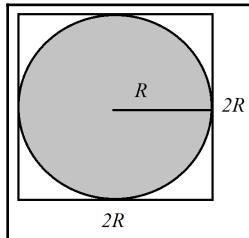
```
In [3]: print s**2
[ 3.00100688e+05  1.00011516e+05  9.27482639e-03  9.26916022e-03
 9.21287015e-03  9.14533995e-03  9.02440213e-03  8.79677106e-03
 8.72804411e-03  8.61862674e-03]
```

```
In [4]: print u[:,0]
[ -7.07105637e-01  7.07107902e-01  -3.11021381e-06  -1.28881106e-06
 -2.04502885e-06  1.54779946e-06  -2.50539756e-06  3.35367849e-06
 -1.68121846e-06  3.36088988e-07]
```

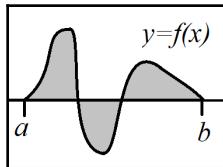
```
In [5]: print u[:,1]
[ -7.07107902e-01  -7.07105637e-01  -5.94599987e-06  1.07432527e-07
 3.17310139e-07  -6.16845739e-07  -9.59202112e-07  -8.96491883e-07
 2.71546742e-06  6.28509406e-06]
```

---

# Chapter 8: The CUDA Device Function Libraries and Thrust



```
In [25]: run monte_carlo_pi.py
Our Monte Carlo estimate of Pi is : 3.14159237769
NumPy's Pi constant is: 3.14159265359
Our estimate passes NumPy's 'allclose' : True
```



```
In [1]: code_string = "%(precision)s x, y; %(precision)s * z;" % {"precision": "double"}
In [2]: code_dict = {'precision' : 'double'}
In [3]: code_double = code_string % code_dict
In [4]: print code_double
double x, y; double * z;
```

```
In [2]: sin_integral = MonteCarloIntegrator()
In [3]: sin_integral.definite_integral()
Out[3]: 2.000000334270522
```

---

```
In [2]: run monte_carlo_integrator.py
The Monte Carlo numerical integration of the function
    f: x -> y =log(x)*_P2(sin(x))
        from x = 11.733 to x = 18.472 is : 8.9999892677
where the expected value is : 8.9999

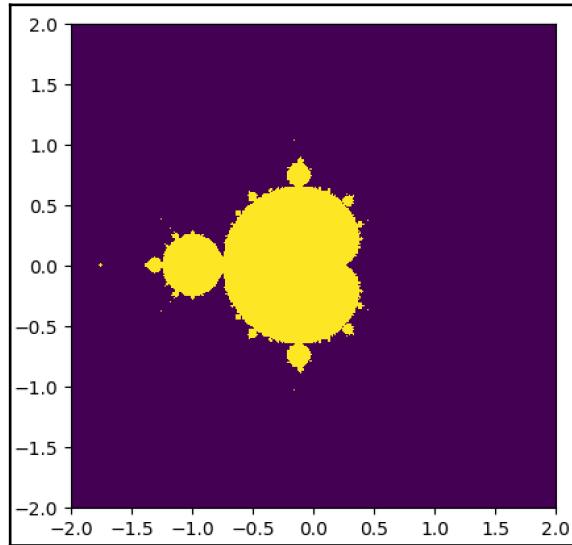
The Monte Carlo numerical integration of the function
    f: x -> y = _R( 1 + sinh(2*x)*_P2(log(x)) )
        from x = 0.9 to x = 4 is : 0.584976671612
where the expected value is : 0.584977

The Monte Carlo numerical integration of the function
    f: x -> y = (cosh(x)*sin(x))/ sqrt( pow(x,3) + _P2(sin(x)))
        from x = 1.85 to x = 4.81 is : -3.34553137474
where the expected value is : -3.34553
```

```
PS C:\Users\btuom\examples\8> .\thrust_dot_product.exe
v[0] == 1
v[1] == 2
v[2] == 3
w[0] == 1
w[1] == 1
w[2] == 1
dot_product(v , w) == 6
```

---

## Chapter 10: Working with Compiled GPU Code



---

# Chapter 11: Performance Optimization in CUDA

```
PS C:\Users\btuom\examples\11> python .\dynamic_hello.py
Hello from thread 0, recursion depth 0!
Hello from thread 1, recursion depth 0!
Hello from thread 2, recursion depth 0!
Hello from thread 3, recursion depth 0!
Launching a new kernel from depth 0 .

-----
Hello from thread 0, recursion depth 1!
Hello from thread 1, recursion depth 1!
Hello from thread 2, recursion depth 1!
Launching a new kernel from depth 1 .

-----
Hello from thread 0, recursion depth 2!
Hello from thread 1, recursion depth 2!
Launching a new kernel from depth 2 .

-----
Hello from thread 0, recursion depth 3!
PS C:\Users\btuom\examples\11>
```

```
PS C:\Users\btuom\examples\11> python .\vectorized_memory.py
Vectorized Memory Test:
First int4: 1, 2, 3, 4
Second int4: 5, 6, 7, 8
First double2: 1.110000, 2.220000
Second double2: 3.330000, 4.440000
```

```
PS C:\Users\btuom\examples\11> python .\atomic.py
Atomic operations test:
add_out: 64
max_out: 127
```

```
PS C:\Users\btuom\examples\11> python .\shfl_xor.py
input array: [ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24
 25 26 27 28 29 30 31]
array after __shfl_xor: [ 1  0  3  2  5  4  7  6  9  8 11 10 13 12 15 14 17 16 19 18 21 20 23 22 25
 24 27 26 29 28 31 30]
```

---

```
PS C:\Users\btruom\examples\11> python .\shfl_sum.py
Input array: [ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24
 25 26 27 28 29 30 31]
Summed value: 496
Does this match with Pythons sum? : True
```

```
PS C:\Users\btruom\examples\11> python .\ptx_assembly.py
x is 123
x is now 0
x is now 1
f is now 3.330000
lane ID: 0
Do split64 / combine64 work? : true
```

```
In [14]: run sum_ker.py
Does sum_ker produces the same value as NumPy's sum (according allclose)? : True
Performing sum_ker / PyCUDA sum timing tests (20 each)...
sum_ker average time duration: 0.00553162831763, PyCUDA's gpuarray.sum average time duration: 0.0278831109579
(Performance improvement of sum_ker over gpuarray.sum: 5.04066964677 )
```